

# HW5\_Ghoshal

Gourav Ghoshal

November 9, 2018

## Problem 1

In this question, we will predict the number of applications received (Apps) using the other variables in the College data set (ISLR package).

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.4.2
```

```
head(College)
```

```
##               Private Apps Accept Enroll Top10perc
## Abilene Christian University    Yes 1660   1232    721      23
## Adelphi University             Yes 2186   1924    512      16
## Adrian College                 Yes 1428   1097    336      22
## Agnes Scott College            Yes  417    349    137      60
## Alaska Pacific University      Yes  193    146     55      16
## Albertson College              Yes  587    479    158      38
##               Top25perc F.Undergrad P.Undergrad Outstate
## Abilene Christian University     52      2885      537    7440
## Adelphi University              29      2683     1227   12280
## Adrian College                  50      1036      99   11250
## Agnes Scott College              89       510      63   12960
## Alaska Pacific University        44       249     869    7560
## Albertson College               62       678      41   13500
##               Room.Board Books Personal PhD Terminal
## Abilene Christian University    3300   450      2200   70      78
## Adelphi University             6450   750      1500   29      30
## Adrian College                 3750   400      1165   53      66
## Agnes Scott College            5450   450       875   92      97
## Alaska Pacific University      4120   800      1500   76      72
## Albertson College              3335   500       675   67      73
##               S.F.Ratio perc.alumni Expend Grad.Rate
## Abilene Christian University   18.1      12    7041     60
## Adelphi University            12.2      16   10527     56
## Adrian College                12.9      30    8735     54
## Agnes Scott College           7.7       37   19016     59
## Alaska Pacific University     11.9       2   10922     15
## Albertson College             9.4       11    9727     55
```

```
names(College)
```

```
## [1] "Private"      "Apps"         "Accept"       "Enroll"       "Top10perc"
## [6] "Top25perc"    "F.Undergrad"  "P.Undergrad"  "Outstate"     "Room.Board"
## [11] "Books"        "Personal"     "PhD"          "Terminal"     "S.F.Ratio"
## [16] "perc.alumni"  "Expend"      "Grad.Rate"
```

```
dim(College)

## [1] 777 18

sum(is.na(College))

## [1] 0
```

(a) Perform best subset selection to the data. What is the best model obtained according to  $C_p$ , BIC and adjusted  $R^2$ ? Show some plots to provide evidence for your answer, and report the coefficients of the best model.

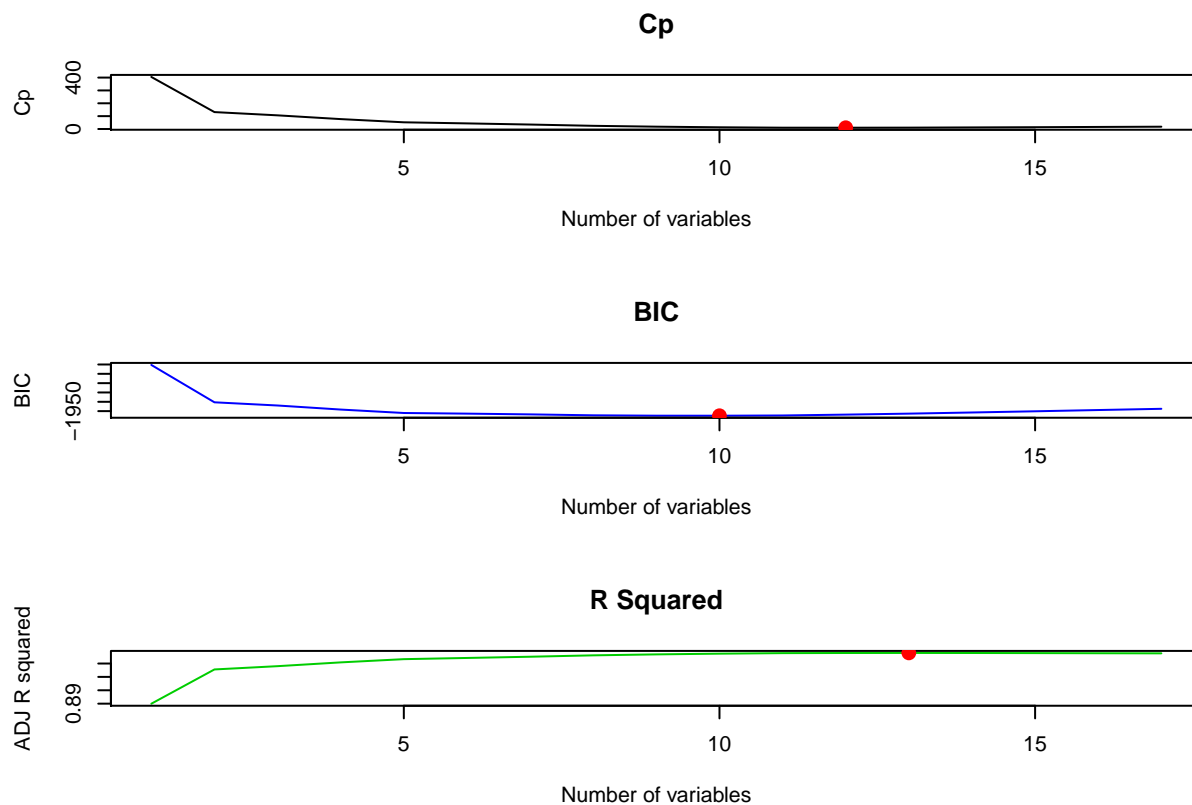
```
# a) Best subset
library(leaps)

## Warning: package 'leaps' was built under R version 3.4.3

attach(College)
best_subset_model = regsubsets(Apps ~., College, nvmax = dim(College)[2]-1)
res_summ = summary(best_subset_model)
names(res_summ)

## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"

par(mfrow=c(3,1))
plot(res_summ$cp,type='l',
     main = 'Cp',xlab = 'Number of variables', ylab = 'Cp')
points(which.min(res_summ$cp), res_summ$cp[which.min(res_summ$cp)],col=2,
       cex=2, pch=20)
plot(res_summ$bic,type='l',col = 4,
     main = 'BIC', xlab = 'Number of variables', ylab = 'BIC')
points(which.min(res_summ$bic), res_summ$bic[which.min(res_summ$bic)],col=2,
       cex=2, pch=20)
plot(res_summ$adjr2,type='l',col = 3,
     main = 'R Squared', xlab = 'Number of variables', ylab = 'ADJ R squared')
points(which.max(res_summ$adjr2), res_summ$adjr2[which.max(res_summ$adjr2)],col=2,
       cex=2, pch=20)
```



```
# Cp ==> 12, BIC==> 10, R-squared ==> 17 (since not adj R-square)
which.min(res_summ$cp)
```

```
## [1] 12
```

```
which.min(res_summ$bic)
```

```
## [1] 10
```

```
which.max(res_summ$adjr2)
```

```
## [1] 13
```

(b) Repeat (a) using forward stepwise selection and backwards stepwise selection. How does your answer compare to the results in (a)?

```
# b) Forward and Backward selection
```

```
## Forward
```

```
f_subset_model = regsubsets(Apps ~., College, nvmax = dim(College)[2]-1,
                             method = 'forward')
```

```
f_summ = summary(f_subset_model)
```

```
names(f_summ)
```

```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
par(mfrow=c(3,1))
```


```
plot(f_summ$cp,type='l',
```

```
main = 'Cp',xlab = 'Number of variables', ylab = 'Cp')
```

```

points(which.min(f_summ$cp), f_summ$cp[which.min(f_summ$cp)],col=2,
       cex=2, pch=20)
plot(f_summ$bic,type='l',col = 4,
     main = 'BIC', xlab = 'Number of variables', ylab = 'BIC')
points(which.min(f_summ$bic), f_summ$bic[which.min(f_summ$bic)],col=2,
       cex=2, pch=20)
plot(f_summ$adjr2,type='l',col = 3,
     main = 'R Squared', xlab = 'Number of variables', ylab = 'ADJ R squared')
points(which.max(f_summ$adjr2), f_summ$adjr2[which.max(f_summ$adjr2)],col=2,
       cex=2, pch=20)

```



```

# Cp ==> 12, BIC==> 10, R-squared ==> 17 (since not adj R-square)
which.min(f_summ$cp)

## [1] 12
which.min(f_summ$bic)

## [1] 10
which.max(f_summ$adjr2)

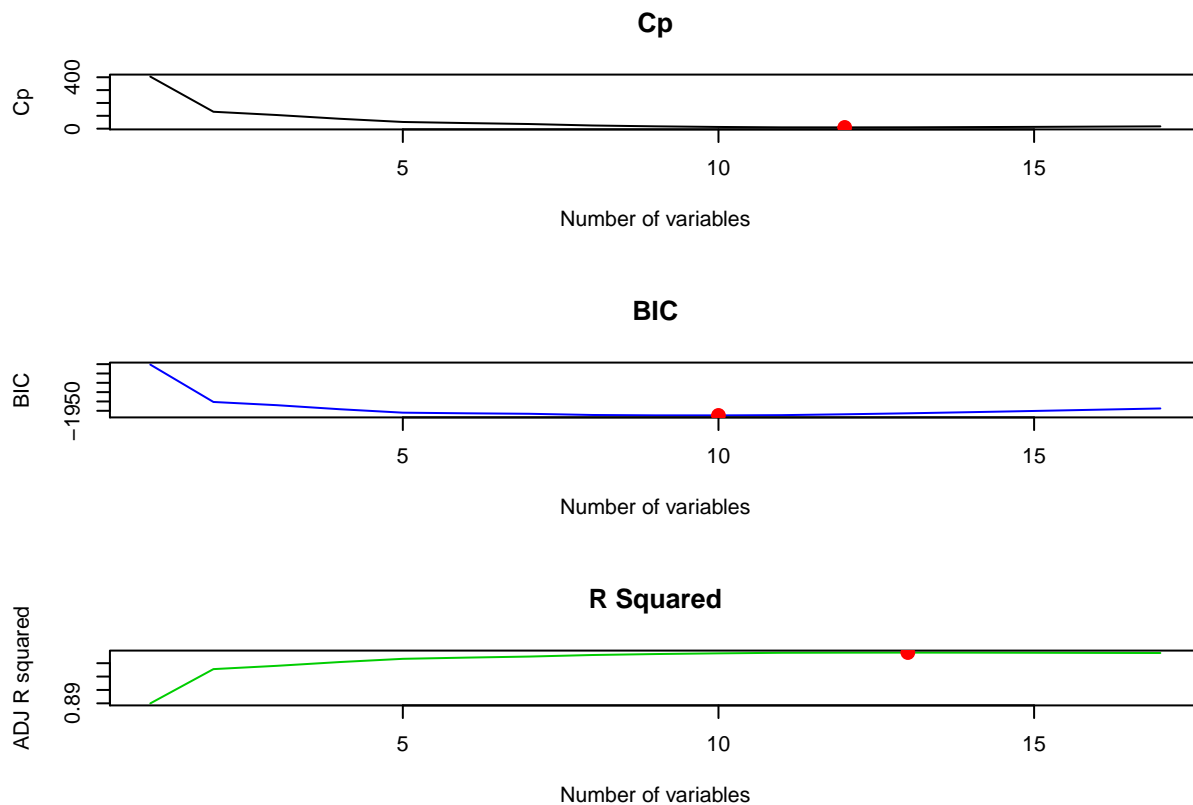
## [1] 13
#backward
b_subset_model = regsubsets(Apps ~., College, nvmax = dim(College)[2]-1,
                           method = 'backward')
b_summ = summary(b_subset_model)

```

```
names(b_summ)
```

```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
par(mfrow=c(3,1))
plot(b_summ$cp,type='l',
     main = 'Cp',xlab = 'Number of variables', ylab = 'Cp')
points(which.min(b_summ$cp), b_summ$cp[which.min(b_summ$cp)],col=2,
       cex=2, pch=20)
plot(b_summ$bic,type='l',col = 4,
     main = 'BIC', xlab = 'Number of variables', ylab = 'BIC')
points(which.min(b_summ$bic), b_summ$bic[which.min(b_summ$bic)],col=2,
       cex=2, pch=20)
plot(b_summ$adjr2,type='l',col = 3,
     main = 'R Squared', xlab = 'Number of variables', ylab = 'ADJ R squared')
points(which.max(b_summ$adjr2), b_summ$adjr2[which.max(b_summ$adjr2)],col=2,
       cex=2, pch=20)
```



```
# Cp ==> 12, BIC==> 10, R-squared ==> 17 (since not adj R-square)
```

```
which.min(b_summ$cp)
```

```
## [1] 12
```

```
which.min(b_summ$bic)
```

```
## [1] 10
```

```
which.max(b_summ$adjr2)
```

```
## [1] 13
```

Part (B) gives the same choices as part (A)

(c) Fit a lasso model on the data. Use cross-validation to select the optimal value of lambda. Create plots of the cross-validation error as a function of lambda and Report the resulting coefficient estimates.

```
#c) Lasso
```

```
x = model.matrix(Apps ~., College)[,-2]
```

```
y = Apps
```

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.4.3
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 3.4.2
```

```
## Loaded glmnet 2.0-13
```


```
grid = 10^seq(10,-2,length = 100)
```

```
model_lasso = glmnet(x,y,alpha = 1,lambda = grid)
```

```
set.seed(42)
```

```
cv_lasso = cv.glmnet(x,y,alpha = 1)
```

```
plot(cv_lasso)
```



```
best_lmnda_lasso = cv_lasso$lambda.min
```


```
predict(model_lasso, type='coefficients', s=best_lmnda_lasso)
```

```
## 18 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1
## (Intercept) -998.38928506
## (Intercept) .
## Accept      1.49980274
## Enroll     -0.26942335
## Top10perc   38.11353044
## Top25perc  -5.67831165
## F.Undergrad .
## P.Undergrad 0.04077278
## Outstate   -0.08946605
## Room.Board 0.11894501
## Books      .
## Personal    0.01100391
## PhD        -4.73850286
## Terminal   -1.31391498
## S.F.Ratio   15.41836974
## perc.alumni -2.04515340
## Expend      0.07584338
## Grad.Rate   5.50890148
```

(d) Fit a ridge regression model on the data. Use cross-validation to select the optimal value of lambda. Create plots of the cross-validation error as a function of lambda. Report the resulting coefficient estimates.

```
#d) Ridge
x = model.matrix(Apps ~., College)[-2]
y = Apps
library(glmnet)
grid = 10^seq(10,-2,length = 100)
model_ridge = glmnet(x,y,alpha = 0,lambda = grid)
set.seed(42)
cv_ridge = cv.glmnet(x,y,alpha = 0)
plot(cv_ridge)
```



```
best_lmlda_ridge = cv_ridge$lambda.min

predict(model_ridge, type='coefficients', s=best_lmlda_ridge)
```

```
## 18 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) -1.984049e+03
## (Intercept) .
## Accept      9.848056e-01
## Enroll      4.912930e-01
## Top10perc   2.480453e+01
## Top25perc   9.576670e-01
```



```
## F.Undergrad 8.599654e-02
## P.Undergrad 3.326895e-02
## Outstate -4.095601e-02
## Room.Board 1.797027e-01
## Books 1.072766e-01
## Personal -3.852404e-03
## PhD -1.948046e+00
## Terminal -2.800965e+00
## S.F.Ratio 2.074290e+01
## perc.alumni -1.045259e+01
## Expend 7.721708e-02
## Grad.Rate 1.047647e+01
```

(e) Now split the data set into a training set and a test set.

```
# e) train- test splitting and model fitting
train_indx = sample(1:nrow(x), nrow(x)/2)
test_indx = (-train_indx)
```

i. Fit the best models obtained in the best subset selection (according to Cp, BIC or adjusted R<sup>2</sup>) to the training set, and report the test error obtained.

```
best_subset_model = regsubsets(Apps ~., College[train_indx,],
                               nvmax = dim(College)[2]-1)
res_summ = summary(best_subset_model)
which.min(res_summ$cp)

## [1] 10

which.min(res_summ$bic)

## [1] 5

which.max(res_summ$adjr2)

## [1] 14

X = model.matrix(Apps~., College[test_indx,])
coef_cp = coef(best_subset_model, id = which.min(res_summ$cp) )
coef_bic = coef(best_subset_model, id = which.min(res_summ$bic) )
coef_adj_rsq = coef(best_subset_model, id = which.min(res_summ$adjr2) )

pred_cp = X[,names(coef_cp)]%*%coef_cp
pred_bic = X[,names(coef_bic)]%*%coef_bic
pred_adj_rsq = X[,names(coef_adj_rsq)]%*%coef_adj_rsq

error_cp = mean((Apps[test_indx]-pred_cp)^2)
error_bic = mean((Apps[test_indx]-pred_bic)^2)
error_adj_rsq = mean((Apps[test_indx]-pred_adj_rsq)^2)

print(error_cp)

## [1] 1562163

print(error_bic)
```

```
## [1] 1599488
print(error_adj_rsqr)
```

```
## [1] 2071866
```

ii. Fit a lasso model to the training set, with lambda chosen by cross validation. Report the test error obtained.

```
best_lmlda_lasso = cv_lasso$lambda.min # 15.07769
pred_lasso = predict(model_lasso, s=best_lmlda_lasso, x)
mean((pred_lasso-y)^2) # 1102253
```

```
## [1] 1102253
```

iii. Fit a ridge regression model to the training set, with lambda chosen by cross validation. Report the test error obtained.

```
best_lmlda_ridge = cv_ridge$lambda.min # 400.4766
pred_ridge = predict(model_ridge, s=best_lmlda_ridge, x)
mean((pred_ridge-y)^2) # 1405339
```

```
## [1] 1405339
```

iv. Compare the test errors obtained in the above analysis (i-iii) and determine the optimal model.

Best model is = Best Subset Selection with Cp

## Problem 2

In the lab, a classification tree was applied to the Carseats data set after converting Sales into a binary response variable. This question will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable (that is, without the conversion).

```
library(ISLR)
head(Carseats)
```

```
##   Sales CompPrice Income Advertising Population Price ShelveLoc Age
## 1  9.50      138     73          11         276    120        Bad  42
## 2 11.22      111     48          16         260     83        Good  65
## 3 10.06      113     35          10         269     80       Medium  59
## 4  7.40      117    100           4         466     97       Medium  55
## 5  4.15      141     64           3         340    128        Bad  38
## 6 10.81      124    113          13         501     72        Bad  78
##   Education Urban  US
## 1         17   Yes Yes
## 2         10   Yes Yes
## 3         12   Yes Yes
## 4         14   Yes Yes
## 5         13   Yes  No
## 6         16   No  Yes
```

```
attach(Carseats)
```

(a) Split the data set into a training set and a test set.

```
train_indx = sample(1:nrow(Carseats), nrow(Carseats)/2)
test_indx = (-train_indx)

train = Carseats[c(train_indx),]
test = Carseats[test_indx,]
```


(b) Fit a regression tree to the training set. Plot the tree, and interpret the results. Then compute the test MSE.

```
library(tree)

## Warning: package 'tree' was built under R version 3.4.3
reg_tree = tree(Sales~., Carseats[train_indx,])
summary(reg_tree)

##
## Regression tree:
## tree(formula = Sales ~ ., data = Carseats[train_indx, ])
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Advertising" "Age" "Price" "CompPrice"
## [6] "Income"
## Number of terminal nodes: 17
## Residual mean deviance: 2.434 = 445.4 / 183
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -5.08200 -0.96750 0.01808 0.00000 1.04100 4.16600

par(mfrow=c(1,1))
plot(reg_tree)
text(reg_tree,pretty = 0)
```



```

pred_tree = predict(reg_tree, Carseats[test_indx,])
mean((pred_tree-Carseats[test_indx,]$Sales)^2)

```


```
## [1] 5.371699
```

(c) Prune the tree obtained in (b). Use cross validation to determine the optimal level of tree complexity. Plot the pruned tree and interpret the results. Compute the test MSE of the pruned tree. Does pruning improve the test error?


```

cv_tree = cv.tree(reg_tree)
plot(cv_tree$size, cv_tree$dev, type='b')

```



```
# There is not much gain after 10  
prune_tree = prune.tree(reg_tree, best = 10)  
plot(prune_tree)  
text(prune_tree)
```



```
prune_pred_tree = predict(prune_tree, Carseats[test_indx,])
mean((prune_pred_tree-Carseats[test_indx,]$Sales)^2)
```

```
## [1] 5.205643
```

```
# Yes, pruned tree produces better test results
```

(d) Use the bagging approach to analyze the data. What test MSE do you obtain? Determine which variables are most important.

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.4.2
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 3.4.2
```

```
library(ipred)
```

```
## Warning: package 'ipred' was built under R version 3.4.2
```

```
bag_tree = bagging(Sales~., Carseats[train_indx,], coob = T)
```

```
pred_bag = predict(bag_tree, Carseats[test_indx,])
```

```
mean((pred_bag-Carseats[test_indx,]$Sales)^2)
```

```
## [1] 3.255091
```

```
varImp(bag_tree)
```

```
##              Overall
## Advertising 1.7237883
## Age         1.6732181
## CompPrice   1.7461440
## Education   0.6413526
## Income      1.2862672
## Population  1.0274530
## Price       2.3206199
## ShelfLoc    1.2002635
## Urban       0.4687182
## US          0.4570515
```

(e) Use random forests to analyze the data. What test MSE do you obtain? Determine which variables are most important.

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.4.4
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

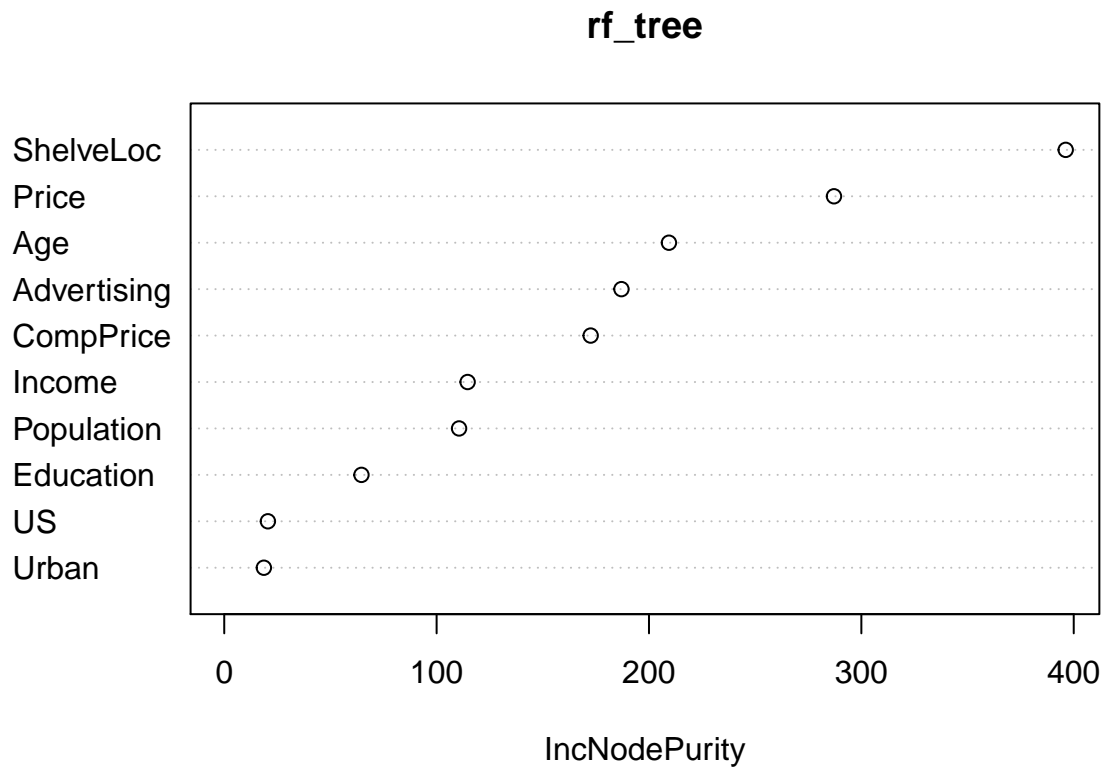
```
rf_tree = randomForest(Sales~., train)
```

```
pred_rf = predict(rf_tree, Carseats[test_indx,])
```

```
mean((pred_rf-Carseats[test_indx,]$Sales)^2)
```

```
## [1] 3.080731
```

```
varImpPlot(rf_tree)
```



ShelveLoc

### Problem 3

In the lab, we applied random forests to the Boston data using `mtry=6` and `ntree=100`.

```
library(MASS)
head(Boston)
```

```
##      crim zn  indus chas   nox    rm  age    dis rad tax ptratio  black
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296    15.3 396.90
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242    17.8 396.90
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242    17.8 392.83
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222    18.7 394.63
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222    18.7 396.90
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222    18.7 394.12
##   lstat medv
## 1   4.98 24.0
## 2   9.14 21.6
## 3   4.03 34.7
## 4   2.94 33.4
## 5   5.33 36.2
## 6   5.21 28.7
```




(a) Consider a more comprehensive range of values for mtry: 1, 2,...,13. Given each value of mtry, find the test error resulting from random forests on the Boston data (using ntree=100). Create a plot displaying the test error rate vs. the value of mtry. Comment on the results in the plot.

```
train_indx = sample(1:nrow(Boston), nrow(Boston)/2)
test_indx = (-train_indx)

attach(Boston)
train = Boston[c(train_indx),]
test = Boston[test_indx,]

library(randomForest)
mse = rep(NA, 13)
for (i in 1:13){
  rf = randomForest(medv ~., train, mtry = i, ntree = 100)
  pred_rf = predict(rf, test)
  mse[i] = mean((pred_rf-test$medv)^2)
}
plot(c(1:13), mse, pch='+', type='l', col=4,
     main = 'Random Forest with different number of predictors',
     xlab = 'Number of Predictors', ylab = 'MSE')
points(c(1:13), mse, pch='+')
```


### Random Forest with different number of predictors



As the number of predictors increases, the error decreases, but there is no improvement with more than 6 predictors

(b) Similarly, consider a range of values for ntree (between 5 to 200). Given each value of ntree, find the test error resulting from random forests (using mtry=6). Create a plot displaying the test error vs. the value of ntree. Comment on the results in the plot.

```
mse = rep(NA, length(c(5:200)))
for (i in 5:200){
  rf = randomForest(medv ~., train, mtry = 6, ntree = i)
  pred_rf = predict(rf, test)
  mse[i] = mean((pred_rf-test$medv)^2)
}
plot(mse, pch='+', type='l', col=4,
     main = 'Random Forest with different number of trees',
     xlab = 'Number of trees', ylab = 'MSE')
```



As the number of trees increases, the error decreases, but there is no improvement after tree size of 100