

Ghoshal_Gourav_HW6

Gourav Ghoshal

November 28, 2018

Problem 1

In this problem, you will use support vector approaches to predict whether a given car gets high or low gas mileage based on the Auto data set in the ISLR package.

(a) Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median. Use this variable as response in the following analysis.

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.4.2
```

```
head(Auto)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1   18         8          307         130   3504          12.0    70     1
## 2   15         8          350         165   3693          11.5    70     1
## 3   18         8          318         150   3436          11.0    70     1
## 4   16         8          304         150   3433          12.0    70     1
## 5   17         8          302         140   3449          10.5    70     1
## 6   15         8          429         198   4341          10.0    70     1
##                                     name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5      ford torino
## 6      ford galaxie 500
```

```
data = Auto
```

```
data$response = ifelse(data$mpg >= median(data$mpg), 1, 0)
```

```
tail(data)
```


```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 392  27         4          151          90   2950          17.3    82     1
## 393  27         4          140          86   2790          15.6    82     1
## 394  44         4           97          52   2130          24.6    82     2
## 395  32         4          135          84   2295          11.6    82     1
## 396  28         4          120          79   2625          18.6    82     1
## 397  31         4          119          82   2720          19.4    82     1
##                                     name response
## 392 chevrolet camaro             1
## 393  ford mustang gl             1
## 394      vw pickup             1
## 395  dodge rampage             1
## 396  ford ranger             1
```

```
## 397      chevy s-10      1
```

(b) Fit a support vector classifier to the data with various values of cost, to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results.

```
# Let's plot the data
attach(data)
plot(mpg, col = (3-response), main = 'Visualize against MPG to verify the response')
```

Visualize against MPG to verify the response



```
# fitting SVM
library(e1071)

## Warning: package 'e1071' was built under R version 3.4.2
data_new = data.frame(x = data[,-c(response,mpg)], y = as.factor(data$response))

set.seed(42)
tune.out=tune(svm ,y~.,data=data_new, kernel ="linear",
              ranges =list(cost=c(0.001 , 0.01, 0.1, 1,5,10,100) ))
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
```

```
##
## - best parameters:
## cost
## 5
##
## - best performance: 0.08435897
##
## - Detailed performance results:
## cost error dispersion
## 1 1e-03 0.13012821 0.05596578
## 2 1e-02 0.08923077 0.04999616
## 3 1e-01 0.09935897 0.05820444
## 4 1e+00 0.08935897 0.05024613
## 5 5e+00 0.08435897 0.05287086
## 6 1e+01 0.08435897 0.05287086
## 7 1e+02 0.08692308 0.05707025

# C= 5 and C = 10, both gave the minimum CV error
# Since higher value of C (cost) gives narrower margin, I decide to select it. cost = 10

## prediction
model = tune.out$best.model
y_hat = predict(model, data_new)

table(predict = y_hat, truth = data_new$y)

##          truth
## predict    0    1
##          0 174  11
##          1  22 185
```

(c) Now repeat (b), this time using SVMs with radial and polynomial kernels, with different values of gamma, degree and cost. Comment on your results.

```
set.seed(1)

tune.out.radial=tune(svm , y~., data=data_new, kernel ="radial",
                    ranges =list(cost=c(0.1 ,1 ,10 ,100 ,1000),
                                degree = c(1,2,3,4,5),
                                gamma=c(0.5,1,2,3,4)))

summary(tune.out.radial)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost degree gamma
##    1      1      1
##
## - best performance: 0.07391026
##
## - Detailed performance results:
```

##	cost	degree	gamma	error	dispersion
## 1	1e-01	1	0.5	0.08923077	0.05559893
## 2	1e+00	1	0.5	0.07391026	0.05315495
## 3	1e+01	1	0.5	0.07897436	0.04567335
## 4	1e+02	1	0.5	0.10448718	0.06862934
## 5	1e+03	1	0.5	0.10448718	0.06072287
## 6	1e-01	2	0.5	0.08923077	0.05559893
## 7	1e+00	2	0.5	0.07391026	0.05315495
## 8	1e+01	2	0.5	0.07897436	0.04567335
## 9	1e+02	2	0.5	0.10448718	0.06862934
## 10	1e+03	2	0.5	0.10448718	0.06072287
## 11	1e-01	3	0.5	0.08923077	0.05559893
## 12	1e+00	3	0.5	0.07391026	0.05315495
## 13	1e+01	3	0.5	0.07897436	0.04567335
## 14	1e+02	3	0.5	0.10448718	0.06862934
## 15	1e+03	3	0.5	0.10448718	0.06072287
## 16	1e-01	4	0.5	0.08923077	0.05559893
## 17	1e+00	4	0.5	0.07391026	0.05315495
## 18	1e+01	4	0.5	0.07897436	0.04567335
## 19	1e+02	4	0.5	0.10448718	0.06862934
## 20	1e+03	4	0.5	0.10448718	0.06072287
## 21	1e-01	5	0.5	0.08923077	0.05559893
## 22	1e+00	5	0.5	0.07391026	0.05315495
## 23	1e+01	5	0.5	0.07897436	0.04567335
## 24	1e+02	5	0.5	0.10448718	0.06862934
## 25	1e+03	5	0.5	0.10448718	0.06072287
## 26	1e-01	1	1.0	0.08923077	0.06408390
## 27	1e+00	1	1.0	0.07391026	0.05176240
## 28	1e+01	1	1.0	0.08916667	0.05799733
## 29	1e+02	1	1.0	0.10698718	0.07386765
## 30	1e+03	1	1.0	0.10698718	0.07386765
## 31	1e-01	2	1.0	0.08923077	0.06408390
## 32	1e+00	2	1.0	0.07391026	0.05176240
## 33	1e+01	2	1.0	0.08916667	0.05799733
## 34	1e+02	2	1.0	0.10698718	0.07386765
## 35	1e+03	2	1.0	0.10698718	0.07386765
## 36	1e-01	3	1.0	0.08923077	0.06408390
## 37	1e+00	3	1.0	0.07391026	0.05176240
## 38	1e+01	3	1.0	0.08916667	0.05799733
## 39	1e+02	3	1.0	0.10698718	0.07386765
## 40	1e+03	3	1.0	0.10698718	0.07386765
## 41	1e-01	4	1.0	0.08923077	0.06408390
## 42	1e+00	4	1.0	0.07391026	0.05176240
## 43	1e+01	4	1.0	0.08916667	0.05799733
## 44	1e+02	4	1.0	0.10698718	0.07386765
## 45	1e+03	4	1.0	0.10698718	0.07386765
## 46	1e-01	5	1.0	0.08923077	0.06408390
## 47	1e+00	5	1.0	0.07391026	0.05176240
## 48	1e+01	5	1.0	0.08916667	0.05799733
## 49	1e+02	5	1.0	0.10698718	0.07386765
## 50	1e+03	5	1.0	0.10698718	0.07386765
## 51	1e-01	1	2.0	0.13512821	0.05778941
## 52	1e+00	1	2.0	0.07634615	0.05879482
## 53	1e+01	1	2.0	0.09423077	0.07220895

## 54	1e+02	1	2.0	0.09923077	0.07799865
## 55	1e+03	1	2.0	0.09923077	0.07799865
## 56	1e-01	2	2.0	0.13512821	0.05778941
## 57	1e+00	2	2.0	0.07634615	0.05879482
## 58	1e+01	2	2.0	0.09423077	0.07220895
## 59	1e+02	2	2.0	0.09923077	0.07799865
## 60	1e+03	2	2.0	0.09923077	0.07799865
## 61	1e-01	3	2.0	0.13512821	0.05778941
## 62	1e+00	3	2.0	0.07634615	0.05879482
## 63	1e+01	3	2.0	0.09423077	0.07220895
## 64	1e+02	3	2.0	0.09923077	0.07799865
## 65	1e+03	3	2.0	0.09923077	0.07799865
## 66	1e-01	4	2.0	0.13512821	0.05778941
## 67	1e+00	4	2.0	0.07634615	0.05879482
## 68	1e+01	4	2.0	0.09423077	0.07220895
## 69	1e+02	4	2.0	0.09923077	0.07799865
## 70	1e+03	4	2.0	0.09923077	0.07799865
## 71	1e-01	5	2.0	0.13512821	0.05778941
## 72	1e+00	5	2.0	0.07634615	0.05879482
## 73	1e+01	5	2.0	0.09423077	0.07220895
## 74	1e+02	5	2.0	0.09923077	0.07799865
## 75	1e+03	5	2.0	0.09923077	0.07799865
## 76	1e-01	1	3.0	0.30128205	0.11716574
## 77	1e+00	1	3.0	0.07634615	0.05867200
## 78	1e+01	1	3.0	0.09673077	0.06968463
## 79	1e+02	1	3.0	0.09923077	0.07316609
## 80	1e+03	1	3.0	0.09923077	0.07316609
## 81	1e-01	2	3.0	0.30128205	0.11716574
## 82	1e+00	2	3.0	0.07634615	0.05867200
## 83	1e+01	2	3.0	0.09673077	0.06968463
## 84	1e+02	2	3.0	0.09923077	0.07316609
## 85	1e+03	2	3.0	0.09923077	0.07316609
## 86	1e-01	3	3.0	0.30128205	0.11716574
## 87	1e+00	3	3.0	0.07634615	0.05867200
## 88	1e+01	3	3.0	0.09673077	0.06968463
## 89	1e+02	3	3.0	0.09923077	0.07316609
## 90	1e+03	3	3.0	0.09923077	0.07316609
## 91	1e-01	4	3.0	0.30128205	0.11716574
## 92	1e+00	4	3.0	0.07634615	0.05867200
## 93	1e+01	4	3.0	0.09673077	0.06968463
## 94	1e+02	4	3.0	0.09923077	0.07316609
## 95	1e+03	4	3.0	0.09923077	0.07316609
## 96	1e-01	5	3.0	0.30128205	0.11716574
## 97	1e+00	5	3.0	0.07634615	0.05867200
## 98	1e+01	5	3.0	0.09673077	0.06968463
## 99	1e+02	5	3.0	0.09923077	0.07316609
## 100	1e+03	5	3.0	0.09923077	0.07316609
## 101	1e-01	1	4.0	0.53820513	0.07278281
## 102	1e+00	1	4.0	0.08660256	0.05385501
## 103	1e+01	1	4.0	0.10179487	0.06661035
## 104	1e+02	1	4.0	0.10179487	0.06966782
## 105	1e+03	1	4.0	0.10179487	0.06966782
## 106	1e-01	2	4.0	0.53820513	0.07278281
## 107	1e+00	2	4.0	0.08660256	0.05385501

```

## 108 1e+01      2  4.0 0.10179487 0.06661035
## 109 1e+02      2  4.0 0.10179487 0.06966782
## 110 1e+03      2  4.0 0.10179487 0.06966782
## 111 1e-01      3  4.0 0.53820513 0.07278281
## 112 1e+00      3  4.0 0.08660256 0.05385501
## 113 1e+01      3  4.0 0.10179487 0.06661035
## 114 1e+02      3  4.0 0.10179487 0.06966782
## 115 1e+03      3  4.0 0.10179487 0.06966782
## 116 1e-01      4  4.0 0.53820513 0.07278281
## 117 1e+00      4  4.0 0.08660256 0.05385501
## 118 1e+01      4  4.0 0.10179487 0.06661035
## 119 1e+02      4  4.0 0.10179487 0.06966782
## 120 1e+03      4  4.0 0.10179487 0.06966782
## 121 1e-01      5  4.0 0.53820513 0.07278281
## 122 1e+00      5  4.0 0.08660256 0.05385501
## 123 1e+01      5  4.0 0.10179487 0.06661035
## 124 1e+02      5  4.0 0.10179487 0.06966782
## 125 1e+03      5  4.0 0.10179487 0.06966782

## Minimum error occur @ cost = 1,degree = 2, gamma = 1

set.seed(1)

tune.out.poly=tune(svm , y~., data=data_new, kernel ="polynomial",
                   ranges =list(cost=c(0.1 ,1 ,10 ,100),
                                degree = c(1,2,3,4,5),
                                gamma=c(0.5,1)))

summary(tune.out.poly)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree gamma
##     1      3    0.5
##
## - best performance: 0.07397436
##
## - Detailed performance results:
##   cost degree gamma      error dispersion
## 1    0.1      1    0.5 0.09185897 0.06069132
## 2    1.0      1    0.5 0.09435897 0.06395767
## 3   10.0      1    0.5 0.08935897 0.05443327
## 4  100.0      1    0.5 0.08679487 0.05567394
## 5    0.1      2    0.5 0.21923077 0.09233341
## 6    1.0      2    0.5 0.19615385 0.07426805
## 7   10.0      2    0.5 0.17083333 0.07329107
## 8  100.0      2    0.5 0.16820513 0.07983348
## 9    0.1      3    0.5 0.07653846 0.05269354
## 10   1.0      3    0.5 0.07397436 0.05464674
## 11  10.0      3    0.5 0.09179487 0.05146406
## 12 100.0      3    0.5 0.08929487 0.04398809
## 13   0.1      4    0.5 0.16064103 0.09334022

```

```
## 14 1.0 4 0.5 0.14525641 0.07506797
## 15 10.0 4 0.5 0.16064103 0.07230196
## 16 100.0 4 0.5 0.18115385 0.08023736
## 17 0.1 5 0.5 0.10461538 0.06888486
## 18 1.0 5 0.5 0.09173077 0.03417795
## 19 10.0 5 0.5 0.10705128 0.05346238
## 20 100.0 5 0.5 0.11474359 0.04855649
## 21 0.1 1 1.0 0.09948718 0.06331482
## 22 1.0 1 1.0 0.09192308 0.05827672
## 23 10.0 1 1.0 0.08935897 0.05443327
## 24 100.0 1 1.0 0.08679487 0.05567394
## 25 0.1 2 1.0 0.18846154 0.07426805
## 26 1.0 2 1.0 0.17576923 0.07691940
## 27 10.0 2 1.0 0.16570513 0.07916097
## 28 100.0 2 1.0 0.17083333 0.08345228
## 29 0.1 3 1.0 0.07397436 0.05464674
## 30 1.0 3 1.0 0.09429487 0.05250366
## 31 10.0 3 1.0 0.09192308 0.04407224
## 32 100.0 3 1.0 0.10724359 0.04971433
## 33 0.1 4 1.0 0.14782051 0.07866510
## 34 1.0 4 1.0 0.16314103 0.07322999
## 35 10.0 4 1.0 0.16339744 0.07437153
## 36 100.0 4 1.0 0.17878205 0.06764193
## 37 0.1 5 1.0 0.09423077 0.05370011
## 38 1.0 5 1.0 0.10692308 0.05300334
## 39 10.0 5 1.0 0.12237179 0.05214415
## 40 100.0 5 1.0 0.12237179 0.05214415
## Minimum error occur @ cost = 0.1, degree = 3, gamma = 1
```

Problem 2

This problem uses the OJ data set in the ISLR package

(a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
library(ISLR)
head(OJ)
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM DiscCH DiscMM SpecialCH
## 1 CH 237 1 1.75 1.99 0.00 0.0 0
## 2 CH 239 1 1.75 1.99 0.00 0.3 0
## 3 CH 245 1 1.86 2.09 0.17 0.0 0
## 4 MM 227 1 1.69 1.69 0.00 0.0 0
## 5 CH 228 7 1.69 1.69 0.00 0.0 0
## 6 CH 230 7 1.69 1.99 0.00 0.0 0
## SpecialMM LoyalCH SalePriceMM SalePriceCH PriceDiff Store7 PctDiscMM
## 1 0 0.500000 1.99 1.75 0.24 No 0.000000
## 2 1 0.600000 1.69 1.75 -0.06 No 0.150754
## 3 0 0.680000 2.09 1.69 0.40 No 0.000000
## 4 0 0.400000 1.69 1.69 0.00 No 0.000000
## 5 0 0.956535 1.69 1.69 0.00 Yes 0.000000
## 6 1 0.965228 1.99 1.69 0.30 Yes 0.000000
```

```
## PctDiscCH ListPriceDiff STORE
## 1 0.000000 0.24 1
## 2 0.000000 0.24 1
## 3 0.091398 0.23 1
## 4 0.000000 0.00 1
## 5 0.000000 0.00 0
## 6 0.000000 0.30 0
```

```
summary(OJ)
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM
## CH:653 Min. :227.0 Min. :1.00 Min. :1.690 Min. :1.690
## MM:417 1st Qu.:240.0 1st Qu.:2.00 1st Qu.:1.790 1st Qu.:1.990
## Median :257.0 Median :3.00 Median :1.860 Median :2.090
## Mean :254.4 Mean :3.96 Mean :1.867 Mean :2.085
## 3rd Qu.:268.0 3rd Qu.:7.00 3rd Qu.:1.990 3rd Qu.:2.180
## Max. :278.0 Max. :7.00 Max. :2.090 Max. :2.290
## DiscCH DiscMM SpecialCH SpecialMM
## Min. :0.00000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.00000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.00000 Median :0.0000 Median :0.0000 Median :0.0000
## Mean :0.05186 Mean :0.1234 Mean :0.1477 Mean :0.1617
## 3rd Qu.:0.00000 3rd Qu.:0.2300 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :0.50000 Max. :0.8000 Max. :1.0000 Max. :1.0000
## LoyalCH SalePriceMM SalePriceCH PriceDiff
## Min. :0.000011 Min. :1.190 Min. :1.390 Min. : -0.6700
## 1st Qu.:0.325257 1st Qu.:1.690 1st Qu.:1.750 1st Qu.: 0.0000
## Median :0.600000 Median :2.090 Median :1.860 Median : 0.2300
## Mean :0.565782 Mean :1.962 Mean :1.816 Mean : 0.1465
## 3rd Qu.:0.850873 3rd Qu.:2.130 3rd Qu.:1.890 3rd Qu.: 0.3200
## Max. :0.999947 Max. :2.290 Max. :2.090 Max. : 0.6400
## Store7 PctDiscMM PctDiscCH ListPriceDiff
## No :714 Min. :0.0000 Min. :0.00000 Min. :0.000
## Yes:356 1st Qu.:0.0000 1st Qu.:0.00000 1st Qu.:0.140
## Median :0.0000 Median :0.00000 Median :0.240
## Mean :0.0593 Mean :0.02731 Mean :0.218
## 3rd Qu.:0.1127 3rd Qu.:0.00000 3rd Qu.:0.300
## Max. :0.4020 Max. :0.25269 Max. :0.440
## STORE
## Min. :0.000
## 1st Qu.:0.000
## Median :2.000
## Mean :1.631
## 3rd Qu.:3.000
## Max. :4.000
```

```
## Creating train-test
set.seed(101)
train_index = sample.int(dim(OJ), size = 800)
train_data = OJ[train_index,]
test_data = OJ[-train_index,]
```


(b) Fit a support vector classifier to the training data using $\text{cost}=0.01$, with Purchase as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.

```
library(e1071)
svm_fit = svm(train_data$Purchase ~., data = train_data,
              kernel = 'linear', cost = 0.01)

summary(svm_fit)

##
## Call:
## svm(formula = train_data$Purchase ~ ., data = train_data, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   0.01
##   gamma:    0.05555556
##
## Number of Support Vectors:  433
##
## ( 218 215 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
## There are 433 support vectors, 218 from CH and 215 from MM
```

(c) What are the training and test error rates?

```
y_hat_train = predict(svm_fit, train_data)
table(predict = y_hat_train, truth = train_data$Purchase)

##          truth
## predict  CH  MM
##      CH 430  77
##      MM  57 236

accuracy_train = (430+236)/(430+77+57+236)
accuracy_train

## [1] 0.8325

y_hat_test = predict(svm_fit, test_data)
table(predict = y_hat_test, truth = test_data$Purchase)

##          truth
## predict  CH  MM
##      CH 148  27
```

```
##      MM 18 77
accuracy_test = (148+77)/(148+27+18+77)
accuracy_test
```

```
## [1] 0.8333333
```

(d) Use the `tune()` function to select an optimal cost. Consider value in the range 0.01 to 10.

```
tune.out=tune(svm ,y~.,data=data_new, kernel ="linear",
              ranges =list(cost=c(0.01, 0.05, 0.1, 0.5, 1, 5, 10)))
summary(tune.out) ## optimal cost = 0.5
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.5
##
## - best performance: 0.07923077
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.08935897 0.03040982
## 2  0.05 0.09705128 0.02682318
## 3  0.10 0.09448718 0.02465186
## 4  0.50 0.07923077 0.03737692
## 5  1.00 0.08179487 0.03180130
## 6  5.00 0.08685897 0.03483004
## 7 10.00 0.08685897 0.03483004

tune.out$best.model

##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = data_new, ranges = list(cost = c(0.01,
##   0.05, 0.1, 0.5, 1, 5, 10)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  0.5
##       gamma: 0.1428571
##
## Number of Support Vectors: 89
```

(e) Compute the training and test error rates using this new value for cost.

```
svm_fit_opt = svm(train_data$Purchase ~., data = train_data,
                  kernel = 'linear', cost = 0.5)
```

```
y_hat_train_opt = predict(svm_fit_opt, train_data)
table(predict=y_hat_train_opt, truth = train_data$Purchase)
```

```
##          truth
## predict  CH  MM
##        CH 429  77
##        MM  58 236
```

```
accuracy_train_opt = (429+236)/(429+77+58+236)
accuracy_train_opt
```

```
## [1] 0.83125
```

```
y_hat_test_opt = predict(svm_fit, test_data)
table(predict=y_hat_test_opt, truth = test_data$Purchase)
```

```
##          truth
## predict  CH  MM
##        CH 148  27
##        MM  18  77
```

```
accuracy_test = (148+77)/(148+27+18+77)
accuracy_test
```

```
## [1] 0.8333333
```

(f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the `tune()` function to select an optimal cost and gamma.

```
# with cost = 0.01
svm_fit = svm(train_data$Purchase ~., data = train_data,
              kernel = 'radial', cost = 0.01)
```

```
summary(svm_fit)
```

```
##
## Call:
## svm(formula = train_data$Purchase ~ ., data = train_data, kernel = "radial",
##      cost = 0.01)
##
##
```

```
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:   0.01
##   gamma:    0.05555556
##
```

```
## Number of Support Vectors:  628
```

```
##
## ( 315 313 )
##
```

```
## Number of Classes:  2
```

```
##
## Levels:
##  CH MM
```

```

## There are 628 support vectors, 315 from CH and 313 from MM class
# Radial Kernel
# Errors using cost = 0.01
y_hat_train = predict(svm_fit, train_data)
table(predict = y_hat_train, truth = train_data$Purchase)

##          truth
## predict  CH  MM
##        CH 487 313
##        MM   0   0

accuracy_train = (487)/(487+0+313+0)
accuracy_train

## [1] 0.60875

y_hat_test = predict(svm_fit, test_data)
table(predict = y_hat_test, truth = test_data$Purchase)

##          truth
## predict  CH  MM
##        CH 166 104
##        MM   0   0

accuracy_test = (166)/(166+0+104+0)
accuracy_test

## [1] 0.6148148

## cv - to find opt params
new_data = data.frame(x = train_data[,-1], y = as.factor(train_data$Purchase))

svm_tune=tune(svm , y~., data=new_data,
              kernel='radial',ranges =list(
                cost=c(0.01 ,0.1 ,1 ,10),gamma=c(0.1,0.5,1,2,3,4)))

summary(svm_tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     1    0.1
##
## - best performance: 0.17375
##
## - Detailed performance results:
##   cost gamma  error dispersion
## 1  0.01   0.1 0.39125 0.06153420
## 2  0.10   0.1 0.18500 0.04706674
## 3  1.00   0.1 0.17375 0.04839436
## 4 10.00   0.1 0.19375 0.04340139
## 5  0.01   0.5 0.39125 0.06153420
## 6  0.10   0.5 0.28500 0.04116363

```

```
## 7 1.00 0.5 0.20000 0.04639804
## 8 10.00 0.5 0.22250 0.05230785
## 9 0.01 1.0 0.39125 0.06153420
## 10 0.10 1.0 0.34375 0.06380580
## 11 1.00 1.0 0.21125 0.05635022
## 12 10.00 1.0 0.22625 0.05382908
## 13 0.01 2.0 0.39125 0.06153420
## 14 0.10 2.0 0.37250 0.06476453
## 15 1.00 2.0 0.22250 0.04594683
## 16 10.00 2.0 0.24875 0.06022239
## 17 0.01 3.0 0.39125 0.06153420
## 18 0.10 3.0 0.38750 0.06152010
## 19 1.00 3.0 0.24000 0.04816061
## 20 10.00 3.0 0.25125 0.05478810
## 21 0.01 4.0 0.39125 0.06153420
## 22 0.10 4.0 0.39125 0.06153420
## 23 1.00 4.0 0.23500 0.05096295
## 24 10.00 4.0 0.25750 0.05374838
```

```
# optimal cost = 1, optimal gamma = 0.1
```

```
## best radial kernel performance -
best_radial = svm(y~., data=new_data, kernel='radial',
                  cost = 1, gamma = 0.1)
y_hat_train = predict(best_radial, new_data)
table(predict=y_hat_train, truth = new_data$y)
```

```
##      truth
## predict CH  MM
##      CH 447  72
##      MM  40 241
```

```
accuracy_train = (447+241)/(447+40+72+241)
accuracy_train # 0.86
```

```
## [1] 0.86
```

```
new_test_data = data.frame(x = test_data[,-1], y = as.factor(test_data$Purchase))
```

```
y_hat_test = predict(best_radial, new_test_data)
table(predict=y_hat_test, truth = new_test_data$y)
```

```
##      truth
## predict CH  MM
##      CH 148  30
##      MM  18  74
```

```
accuracy_test = (148+74)/(148+18+30+74)
accuracy_test # 0.822
```

```
## [1] 0.8222222
```

(g) Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set degree=2. Use the tune() function to select an optimal cost.

```

# Polynomial kernel
# with cost = 0.01
svm_fit = svm(train_data$Purchase ~., data = train_data,
               kernel = 'polynomial', degree = 2)

summary(svm_fit)

##
## Call:
## svm(formula = train_data$Purchase ~ ., data = train_data, kernel = "polynomial",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   1
##   degree:   2
##   gamma:   0.05555556
##   coef.0:   0
##
## Number of Support Vectors:  451
##
## ( 227 224 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
##
## There are 451 support vectors, 227 from CH and 224 from MM class

# Errors using degree = 2
y_hat_train = predict(svm_fit, train_data)
table(predict=y_hat_train, truth = train_data$Purchase)

##           truth
## predict  CH  MM
##      CH 459 112
##      MM  28 201

accuracy_train = (459+201)/(459+112+28+201)
accuracy_train

## [1] 0.825

y_hat_test = predict(svm_fit, test_data)
table(predict=y_hat_test, truth = test_data$Purchase)

##           truth
## predict  CH  MM
##      CH 153  41
##      MM  13  63

```

```

accuracy_test = (153+63)/(153+41+13+63)
accuracy_test

## [1] 0.8
new_data = data.frame(x = train_data[,-1], y = as.factor(train_data$Purchase))

svm_tune=tune(svm , y~., data=new_data,
              kernel="polynomial",ranges =list(cost=c(0.01, 0.1 ,1 ,10
              ,100, 1000),
              degree = 2))
summary(svm_tune)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##   1000      2
##
## - best performance: 0.17125
##
## - Detailed performance results:
##   cost degree   error dispersion
## 1 1e-02      2 0.39000 0.03670453
## 2 1e-01      2 0.32000 0.05688683
## 3 1e+00      2 0.19375 0.06802012
## 4 1e+01      2 0.18500 0.04362084
## 5 1e+02      2 0.17375 0.04185375
## 6 1e+03      2 0.17125 0.03230175

# optimal cost = 100

## best polynomial kernel performance -
best_poly = svm(y~., data=new_data, kernel='polynomial', degree = 2,
               cost = 100)
y_hat_train = predict(best_poly, new_data)
table(predict =y_hat_train, truth = new_data$y)

##           truth
## predict  CH  MM
##      CH 442  57
##      MM  45 256

accuracy_train = (442+256)/(442+45+57+256)
accuracy_train # 0.8725

## [1] 0.8725
new_test_data = data.frame(x = test_data[,-1], y = as.factor(test_data$Purchase))

y_hat_test = predict(best_poly, new_test_data)
table(predict =y_hat_test, truth = new_test_data$y)

##           truth

```

```
## predict CH MM
##      CH 147 26
##      MM  19 78
accuracy_test = (147+78)/(147+19+26+78)
accuracy_test # 0.833
## [1] 0.8333333
```

(h) Overall, which approach seems to give the best results on this data?

```
## Overall the Polynomial kernel performs better as both the training and testing accuracy is higher
```